

YUANJUN (JOHNNY) DAI

216-269-2394 • yxd429@case.edu • Cleveland, OH
[GitHub](#) • [Portfolio](#) • [Blog](#) • [LinkedIn](#)

SUMMARY

AI Infrastructure and Applied AI Engineer, PhD candidate at CWRU. Specializes in the security, optimization, and observability of ML/LLM training and inference infrastructure. 8 research papers (7 first-author) at IEEE, ACM, and Springer; **2+ years of systems software engineering** at Cisco and Broadcom; 2+ years of **cross-functional** product ownership at Midea and Xiaomi.

EDUCATION

Ph.D. Candidate in Computer Science (AI Infrastructure) 08/2019 – Sep. 2026 (exp.)
Case Western Reserve University

M.S. in Electrical and Computer Engineering (Computer Architecture)
University of Pittsburgh

B.S. in Electrical Engineering
Northeastern University

SKILLS

Languages	Python, C/C++
ML Frameworks	PyTorch (DDP/FSDP), vLLM, Megatron-LM, NCCL, BytePS, scikit-learn
Parallelism Paradigms	Data Parallel (DDP/FSDP), Tensor Parallel (TP), Pipeline Parallel (PP), Hybrid Parallel
Infrastructure	Kubernetes, Docker, CI/CD Pipelines (GitHub Actions), ArgoCD, Slurm
Systems & APIs	FastAPI, REST API, gRPC
Observability	eBPF (Kernel Tracing), Prometheus, Grafana, DCGM Exporter, Intel RAPL
Cloud & HPC	Lambda Labs HPC, Ohio Supercomputer Center, NSF FABRIC, CloudLab
Networking	SDN (Software-Defined Networking), P4, NPU/ASIC (Broadcom BCM SDK)

AI INFRASTRUCTURE & SYSTEMS EXPERIENCE

Research Assistant, AI Infrastructure & ML Systems 08/2019 – Sep. 2026 (exp.)
Case Western Reserve University

Research focus: **security, optimization, and observability** of **ML/LLM** training and inference infrastructure.

DYNAMIX – Reinforcement Learning-Based Adaptive Batch Scheduling for Distributed ML | [Paper](#) / [Code](#)

- Designed a **PPO-based Reinforcement Learning (RL)** scheduler monitoring hardware signals (CPU/memory, bandwidth, retransmission rate via **eBPF** probes) and statistical signals (loss convergence, generalization) to dynamically adapt batch size
- Co-scales learning rate via **Linear Scaling Rule**; evaluated on **PyTorch DDP** and **Parameter Server (BytePS)**
- Achieves **46% wall-clock speedup** and **~6% accuracy gain** across **8–64 A100 GPUs** over fixed batch size baseline

AI-Powered Job Hunting Agent | [Code](#)

- Built a production-oriented **multi-agent AI platform** for automated job discovery, semantic matching, and resume optimization across LinkedIn, Indeed, Dice, YC Jobs, and HackerNews
- Designed a **LangGraph-based orchestration pipeline** with specialized agents for collection, matching, refinement, trend analysis, and validation
- Developed a **6-dimension LLM-based job matching framework** via careful **prompt engineering** evaluating skills, experience, seniority, authorization, and company preference
- Implemented a **capability-aware multi-model routing pipeline** for analytical planning and recruiter-oriented resume optimization
- Engineered **FastAPI** backend with cross-platform deduplication, health monitoring endpoints, and deployment-ready automation

Domain-Adapted LLMs for Clinical Risk Adjustment Coding | [Paper](#)

- Built a 115-label HCC multi-label classifier on highly imbalanced EHR data (MIMIC-V), improving long-tail performance via **inverse-frequency weighted binary cross-entropy loss** and **per-label threshold calibration**, achieving **0.775 macro-F1 / 0.742 micro-F1 / 0.9362 macro-AUC**.
- Trained **PubMedBERT** using **PyTorch DDP** on 2×NVIDIA RTX A6000 GPUs for efficient scaling, while fine-tuning **Qwen-7B** using distribution-aware **few-shot prompting** (5-shot with negative, frequent, and rare-label demonstrations) to **mitigate hallucination** and enforce structured HCC outputs with improved sensitivity to sparse labels; scaled full-parameter training via **NeMo + Megatron-LM** (TP=1, PP=1) with **PyTorch FSDP (ZeRO-3)**, **FlashAttention**, and **activation checkpointing**, enabling training under GPU memory constraints (avoiding OOM).

Scalable LLM Serving Infrastructure on Kubernetes | [Code](#) / [Video](#)

- Independently designed a self-hosted **K8s** LLM inference platform (NGINX → **FastAPI** → **vLLM**) with **GitOps CI/CD** (**ArgoCD** + GitHub Actions) across 50+ deployments; validated on **RTX 4050** (**GPU time-slicing**, **Horizontal Pod Autoscaler (HPA)** 1↔2) and **Lambda Labs A100** (**Multi-Instance GPU (MIG)** 7×1g.5gb, **HPA** 1↔7) with **91.3% peak GPU compute** under extreme stress
- Built 4-layer GPU observability (**DCGM Exporter** + **Prometheus** + **Grafana**: GPU hardware, vLLM metrics, FastAPI gateway, K8s cluster)

PRISM – Priority-Based Selective Reliability for Distributed ML Training | [Paper](#)

- Designed a transport-layer library (**NCCL/Gloo** → **UDP**) hooking into **PyTorch DDP**'s AllReduce at bucket-ready events; classifies gradients by magnitude into high-priority (full **SACK**-based retransmission) and low-priority (local compensation via cached values)
- Supports kernel-bypass (**LibVMA**) and kernel-based (**GSO/GRO**) variants; achieves **8–15% throughput gain for CNNs** and **15–21% for LLMs** over TCP, matching NDP latency at 144-worker scale

In-Kernel eBPF Sketch Systems for Network and System Resource Monitoring | [Paper \(PSketch\)](#) / [Code](#) · [Paper \(eHashPipe\)](#) / [Code](#)

- PSketch**: first in-kernel priority-aware sketch using **eBPF**; high-priority flows tracked via in-kernel hashmap ($O(1)$), low-priority via **3-layer Count-Min Sketch**; enables **$O(n)$ Top-K detection** with sub-linear memory; achieves **96% Top-K accuracy**, **96.4% retransmission recall**, <1.1% throughput overhead at 10 Gbps
- eHashPipe**: dual-pipeline **eBPF** framework with **$O(n)$ Top-K tracking** — (1) HashPipe-based memory profiler with cascading slot eviction and ptr2size deallocation tracking, (2) per-thread CPU tracker via sched_switch tracepoint; delivers **10–14× finer temporal resolution** than top/htop, ~11ms latency

Training-Time Side-Channel Attacks on Distributed ML Systems | [Paper](#) / [Code](#)

- Analyzed DML computation/communication behavior** in **MXNet BSP** Parameter Server: instrumented push/pull primitives to profile per-layer network traffic patterns; modified MXNet source code for ground-truth collection to capture **DML network behavior** at layer granularity
- Profiled **ML energy consumption** via **Intel RAPL** (PP0+DRAM, 50μs via powercap/sysfs); applied **sliding window** algorithms to isolate forward passes; used **band-pass filtering** and **K-means clustering** for noise reduction and layer segmentation; applied **polynomial regression** for per-layer and **activation function energy profiling**; classified layer types via **Decision Tree** and recovered Conv/FC dimensions via weighted-average **MLP ensemble** — exposing DNN architecture stealing surfaces in MLaaS environments

INDUSTRY EXPERIENCE

Software Engineer II

06/2015 – 08/2016

Cisco Systems, Inc.

- Developed NPU/ASIC features for a 100 Gbps core router (Azure, AWS, Alibaba Cloud) using Broadcom BCM SDK.
- Contributed to QoS and ECMP modules, reducing forwarding latency under heavy load.

Software Engineer Intern

10/2014 – 05/2015

Broadcom (Emulex)

- Re-architected RTOS data acquisition with DMA + circular-buffer, **reducing CPU load by 30%**; ported BladeEngine ASIC management CLI from Linux to Windows Server.

Product Manager

09/2016 – 06/2018

Midea Group

- Led technical PM for **Fun Oven II** across ML model training, Docker deployment on Alibaba Cloud, mobile app, and culinary algorithm design; won **Red Dot Design Award**, debuted at 2018 AWE Show · [Red Dot ↗](#) · [PR Newswire ↗](#)

Technical Product Manager

07/2018 – 01/2019

Xiaomi

- Defined IoT appliance product requirements via competitive analysis and user research; assessed technical feasibility and delivered architecture proposals to align engineering and business stakeholders.

PUBLICATIONS

8 papers total, 7 first-author, 5 published at IEEE CCNC, IEEE ICNC, ACM TOIT, and Springer SecureComm; 3 under review (incl. ACM TOPS). Full list: johnny-dai-git.github.io/portfolio

PATENTS

Oven Temperature Control Method, Device and Computer-Readable Storage Medium — Granted CN 108594898 B | Issued July 23, 2021 | Assignee: Midea Group Co., Ltd. | Dai Yuanjun et al.