

# YUANJUN (JOHNNY) DAI

216-269-2394 • yxd429@case.edu • Cleveland, OH

[GitHub](#) • [Portfolio](#) • [Blog](#) • [LinkedIn](#)

## SUMMARY

---

AI Systems and Applied AI Engineer, PhD candidate at CWRU. Specializes in the security, optimization, and observability of ML/LLM training and inference infrastructure. 8 research papers (7 first-author) at IEEE, ACM, and Springer; 2+ years of systems software engineering at Cisco and Broadcom; 2+ years of cross-functional product ownership at Midea and Xiaomi.

## EDUCATION

---

### Ph.D. Candidate in Computer Science (AI Infrastructure)

Case Western Reserve University

08/2019 – Sep. 2026 (exp.)

### M.S. in Electrical and Computer Engineering (Computer Microarchitecture)

University of Pittsburgh

### B.S. in Electrical Engineering

Northeastern University

## RESEARCH & SYSTEMS ENGINEERING EXPERIENCE

---

### Research Assistant, AI Infrastructure & ML Systems

Case Western Reserve University

- **DYNAMIX – RL-Based Adaptive Batch Scheduling for Distributed ML** | [Paper](#) / [Code](#)  
PPO-based centralized RL scheduler dynamically adapting batch size via real-time system signals; evaluated on PS (BytePS) and Ring AllReduce (PyTorch DDP) architectures; achieves **46% training speedup** over optimal fixed batch size strategy with improved model generalization, on up to **64 A100 GPUs** at Ohio Supercomputer Center and Lambda Labs HPC.
- **PRISM – Priority-Aware Transport for Distributed ML Training** | [Paper \(Under Review\)](#)  
Priority-based selective reliability transport between NCCL/Gloo and UDP; magnitude-based gradient classification with full retransmission for high-priority and local compensation for low-priority; dual congestion windows, adaptive window borrowing, GPU-accelerated TopK, SACK-based tracking, kernel-bypass (LibVMA) and GSO/GRO optimizations; achieves **8–15% throughput gain for CNNs and 15–21% for LLMs** over TCP on 32 GPUs (OSC) and cross-site FABRIC testbed; matches NDP latency at 144-worker scale on commodity infrastructure.
- **Training-Time Side-Channel Attacks on Distributed ML Systems** | [Paper](#) / [Code](#)  
Analyzed MXNet BSP PS-architecture communication patterns via push/pull instrumentation; modified MXNet source code for ground-truth collection (framework instrumentation); cross-correlated network traffic and Intel RAPL power side channels with outlier filtering and polynomial regression for energy profiling; achieved **>99% layer-type accuracy and <1.5% tensor-size error** across ZFNet, AlexNet, and VGG variants — exposing DNN architecture stealing surfaces in MLaaS co-location environments.
- **Priority-Based eBPF Network Flow Telemetry for AI / DML Infrastructure** | [Paper](#) / [Code](#)  
In-kernel eBPF priority-aware sketch framework (priority table + heavy flow table + 3-layer CMS); sequence-number-based retransmission detection; atomic operations for thread-safe concurrent updates; voting-based collision resolution and CMS minimum-value flow reconstruction; achieved **96% Top-K accuracy, 100% priority-packet recall**, and **<1.1% throughput overhead** at 10 Gbps on FABRIC testbed with CAIDA 2019 traffic.
- **High-Resolution eBPF System Resource Profiling for DML Workloads** | [Paper](#) / [Code](#)  
Developed a per-PID eBPF profiler with in-kernel Top-K tracking, delivering **10–14× finer temporal resolution** than top/htop, ~11 ms latency, and **>90% Top-K fidelity** on commodity Linux servers.
- **Scotch – Elastic SDN Overlay for Control-Plane Scalability** | [Paper](#) / [Code](#)  
Designed a distributed control-plane overlay using OvS-based traffic offloading and large-flow migration; **reduced controller congestion by >70%** on a 34-node GENI Clos topology.
- **Enterprise-Grade LLM Production Deployment Platform on Kubernetes** | [Code](#)  
Built a K8s microservices platform (Gateway / Router / Worker) supporting vLLM and TensorRT-LLM inference with ArgoCD GitOps, Prometheus/Grafana observability, and CI/CD automation across **50+ Kubernetes deployments**.
- **Domain-Adapted LLMs for Clinical Risk Adjustment Coding** | [Paper](#)  
Fine-tuned PubMedBERT and Qwen-7B for HCC coding from EHR notes; achieved **macro-/micro-F1 of 0.775/0.742 and 0.81/0.802** on MIMIC-IV with imbalance-aware training and per-label calibration.
- **BiLSTM-Based Side-Channel Modeling for Distributed ML Analysis** | [Paper](#)  
Designed a BiLSTM pipeline aligning CPU/memory energy traces with per-layer timestamps; reliably distinguishes Conv, FC, and activation layers during distributed ML training using ensemble learning.

INDUSTRY EXPERIENCE

Technical Product Manager

Xiaomi

- Market-facing PM role: conducted competitive analysis and user research to define product requirements for a smart IoT appliance.
- Assessed technical feasibility and complexity of proposed features; delivered initial technical proposals covering architecture trade-offs and implementation risks to align engineering and business stakeholders.

07/2018 – 01/2019

Product Manager

Midea Group

- Served as technical PM for **Fun Oven II**, a multi-disciplinary smart appliance spanning data collection, ML model training & deployment, mobile app integration, cloud backend, and culinary algorithm design.
- Participated in resolving technical challenges across all dimensions — from training AI image recognition models and deploying Docker services on Alibaba Cloud, to defining cooking curves and app UX.
- Product debuted at **2018 AWE Show** and won the **Red Dot Design Award** · [Red Dot ↗](#) · [PR Newswire ↗](#)

09/2016 – 06/2018

Software Engineer II

Cisco Systems, Inc.

- Developed NPU/ASIC features for a 100 Gbps core router (Azure, AWS, Alibaba Cloud) using Broadcom BCM SDK.
- Contributed to QoS and ECMP modules, reducing forwarding latency under heavy load.

06/2015 – 08/2016

Software Engineer Intern

Broadcom (Emulex)

- Replaced polling-based data acquisition with a DMA + circular-buffer architecture in an RTOS pipeline, **reducing CPU load by 30%** and improving sampling rate and control-loop responsiveness.
- Ported a Linux-only management CLI for BladeEngine ASIC-based CNA hardware to Windows Server, adapting system calls, threading models, and error handling.

10/2014 – 05/2015

Engineer Intern

Tenova I2S

- Re-architected an RTOS data acquisition pipeline using DMA and circular buffers, **reducing CPU load by 42%** and improving sampling rate and control-loop responsiveness.

01/2014 – 04/2014

ADDITIONAL SYSTEMS ENGINEERING EXPERIENCE

16-bit MIPS-Like Pipelined CPU Design and Implementation

· [Report ↗](#)

- Designed an end-to-end pipelined CPU from custom ISA through RTL, synthesis, and layout; 5-stage pipeline with hazard handling and branch prediction, achieving **~3× speedup** over single-cycle design.

PUBLICATIONS

8 papers total, 7 first-author, 5 published at IEEE CCNC, IEEE ICNC, ACM TOIT, and Springer SecureComm; 3 under review (incl. ACM TOPS). Full list: [johnny-dai-git.github.io/portfolio](https://johnny-dai-git.github.io/portfolio)

SKILLS

|                |                                                                                              |
|----------------|----------------------------------------------------------------------------------------------|
| Languages      | Python, C/C++, Java, Bash                                                                    |
| AI / ML        | PyTorch (DDP/FSDP), TensorFlow, BytePS, Megatron-LM, NCCL, DeepSpeed, vLLM, TensorRT-LLM     |
| Infrastructure | Kubernetes, Docker, CI/CD, ArgoCD, Helm, Slurm, HPC, GitHub Actions                          |
| Observability  | eBPF, Prometheus, Grafana, DCGM Exporter, Intel RAPL, Kernel Tracing                         |
| Cloud & HPC    | Lambda Labs HPC, NSF FABRIC, CloudLab, SDN (OpenFlow / Ryu), P4, NPU/ASIC (Broadcom BCM SDK) |
| Research       | Distributed ML Systems, ML Security, Side-Channel Analysis, RL-based Scheduling              |

PATENTS

CN 108594898 B — Oven Temperature Control Method, Device and Computer-Readable Storage Medium. Issued: July 23, 2021. Assignee: Midea Group Co., Ltd. Inventors: Shao Guangyuan, Dai Yuanjun, Zhu Jiele, Cheng Gang, Gong Jun.